

# House Loan Data Analysis

## Deep Learning Project\_01

```
In [1]: # Lets first import the important library.
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # Let ignore the un necessary warnings
import warnings
warnings.filterwarnings("ignore")
```

### Load the dataset that is given to you

```
In [3]: df = pd.read_csv('loan_data.csv')
```

```
In [4]: df.head()
```

```
Out[4]:
```

	SK_ID_CURR	TARGET	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN
0	100002	1	Cash loans	M	N	Y	
1	100003	0	Cash loans	F	N	N	
2	100004	0	Revolving loans	M	Y	Y	
3	100006	0	Cash loans	F	N	Y	
4	100007	0	Cash loans	M	N	Y	

5 rows × 122 columns

```
In [5]: df.shape
```

```
Out[5]: (307511, 122)
```

### Check for null values in the dataset

```
In [6]: pd.set_option('display.max_rows', 500)
```

```
In [7]: ((df.isnull().sum()/df.shape[0])*100)
```

```
Out[7]: SK_ID_CURR          0.000000
TARGET          0.000000
NAME_CONTRACT_TYPE 0.000000
CODE_GENDER      0.000000
FLAG_OWN_CAR     0.000000
FLAG_OWN_REALTY  0.000000
CNT_CHILDREN     0.000000
AMT_INCOME_TOTAL 0.000000
AMT_CREDIT       0.000000
AMT_ANNUITY      0.003902
AMT_GOODS_PRICE  0.090403
NAME_TYPE_SUITE  0.420148
NAME_INCOME_TYPE 0.000000
NAME_EDUCATION_TYPE 0.000000
NAME_FAMILY_STATUS 0.000000
NAME_HOUSING_TYPE 0.000000
REGION_POPULATION_RELATIVE 0.000000
DAYS_BIRTH       0.000000
DAYS_EMPLOYED    0.000000
DAYS_REGISTRATION 0.000000
```

```
In [8]: # Lets drop the columns in which max of the data point is missed.
loan_df = df.drop(['OWN_CAR_AGE', 'EXT_SOURCE_1', 'APARTMENTS_AVG', 'BASEMENTAREA_AVG', 'YEARS_BI
                'YEARS_BUILD_AVG', 'COMMONAREA_AVG', 'ELEVATORS_AVG', 'ENTRANCES_AVG', 'FLOORSMA
                'LANDAREA_AVG', 'LIVINGAPARTMENTS_AVG', 'LIVINGAREA_AVG', 'NONLIVINGAPARTMENTS_AV
                'APARTMENTS_MODE', 'BASEMENTAREA_MODE', 'YEARS_BEGINEXPLUATATION_MODE', 'YEARS_BI
                'ELEVATORS_MODE', 'ENTRANCES_MODE', 'FLOORSMAX_MODE', 'FLOORSMIN_MODE', 'LANDAREA
                'LIVINGAREA_MODE', 'NONLIVINGAPARTMENTS_MODE', 'NONLIVINGAREA_MODE', 'APARTMENTS
                'YEARS_BEGINEXPLUATATION_MEDI', 'YEARS_BUILD_MEDI', 'COMMONAREA_MEDI', 'ELEVATORS
                'FLOORSMAX_MEDI', 'FLOORSMIN_MEDI', 'LANDAREA_MEDI', 'LIVINGAPARTMENTS_MEDI', 'LI
                'NONLIVINGAPARTMENTS_MEDI', 'NONLIVINGAREA_MEDI', 'FONDKAPREMONT_MODE', 'HOUSETY
                'WALLSMATERIAL_MODE', 'EMERGENCYSTATE_MODE', 'SK_ID_CURR'], axis=1)
```

```
In [9]: loan_df = loan_df.dropna(subset=['OCCUPATION_TYPE', 'AMT_REQ_CREDIT_BUREAU_HOUR', 'EXT_SOURCE_3',
                'EXT_SOURCE_2', 'NAME_TYPE_SUITE', 'AMT_ANNUITY'])
```

```
In [10]: loan_df.isnull().sum()
```

```
Out[10]: TARGET                                0
NAME_CONTRACT_TYPE                            0
CODE_GENDER                                   0
FLAG_OWN_CAR                                  0
FLAG_OWN_REALTY                              0
CNT_CHILDREN                                 0
AMT_INCOME_TOTAL                             0
AMT_CREDIT                                    0
AMT_ANNUITY                                  0
AMT_GOODS_PRICE                              0
NAME_TYPE_SUITE                              0
NAME_INCOME_TYPE                             0
NAME_EDUCATION_TYPE                         0
NAME_FAMILY_STATUS                           0
NAME_HOUSING_TYPE                           0
REGION_POPULATION_RELATIVE                   0
DAYS_BIRTH                                   0
DAYS_EMPLOYED                               0
DAYS_REGISTRATION                           0
DAYS_TO_REPAYMENT                           0
```

```
In [11]: loan_df['NAME_CONTRACT_TYPE'].value_counts()
```

```
Out[11]: Cash loans          151480
Revolving loans           16252
Name: NAME_CONTRACT_TYPE, dtype: int64
```

```
In [12]: loan_df['CODE_GENDER'].value_counts()
```

```
Out[12]: F          103737
M           63993
XNA              2
Name: CODE_GENDER, dtype: int64
```

```
In [13]: loan_df['FLAG_OWN_CAR'].value_counts()
```

```
Out[13]: N          102179
Y           65553
Name: FLAG_OWN_CAR, dtype: int64
```

```
In [14]: loan_df['FLAG_OWN_REALTY'].value_counts()
```

```
Out[14]: Y          113749
N           53983
Name: FLAG_OWN_REALTY, dtype: int64
```

```
In [15]: loan_df['NAME_TYPE_SUITE'].value_counts()
```

```
Out[15]: Unaccompanied      137289
         Family             21103
         Spouse, partner    6312
         Children          1502
         Other_B            928
         Other_A            455
         Group of people    143
         Name: NAME_TYPE_SUITE, dtype: int64
```

```
In [16]: loan_df['NAME_INCOME_TYPE'].value_counts()
```

```
Out[16]: Working            106020
         Commercial associate 46553
         State servant       15144
         Student             10
         Businessman         3
         Pensioner           2
         Name: NAME_INCOME_TYPE, dtype: int64
```

```
In [17]: loan_df['NAME_EDUCATION_TYPE'].value_counts()
```

```
Out[17]: Secondary / secondary special 115722
         Higher education               44477
         Incomplete higher              6010
         Lower secondary                1433
         Academic degree                90
         Name: NAME_EDUCATION_TYPE, dtype: int64
```

```
In [18]: loan_df['NAME_FAMILY_STATUS'].value_counts()
```

```
Out[18]: Married           111844
         Single / not married 24171
         Civil marriage      16515
         Separated           10769
         Widow              4433
         Name: NAME_FAMILY_STATUS, dtype: int64
```

```
In [19]: loan_df['NAME_HOUSING_TYPE'].value_counts()
```

```
Out[19]: House / apartment    147526
         With parents          9172
         Municipal apartment   5916
         Rented apartment     2909
         Office apartment     1517
         Co-op apartment      692
         Name: NAME_HOUSING_TYPE, dtype: int64
```

```
In [20]: loan_df['OCCUPATION_TYPE'].value_counts()
```

```
Out[20]: Laborers          43437
         Sales staff       24254
         Core staff        22669
         Managers          17787
         Drivers           14546
         High skill tech staff 9264
         Accountants        8135
         Medicine staff     7223
         Security staff     5218
         Cooking staff      4633
         Cleaning staff     3637
         Private service staff 1975
         Low-skill Laborers  1443
         Secretaries        1047
         Waiters/barmen staff 1010
         Realty agents       569
         HR staff           460
         IT staff           425
         Name: OCCUPATION_TYPE, dtype: int64
```

```
In [21]: loan_df['WEEKDAY_APPR_PROCESS_START'].value_counts()
```

```
Out[21]: WEDNESDAY      28655
          TUESDAY      28469
          THURSDAY     27600
          FRIDAY       27536
          MONDAY       27298
          SATURDAY     19159
          SUNDAY       9015
          Name: WEEKDAY_APPR_PROCESS_START, dtype: int64
```

```
In [22]: loan_df['ORGANIZATION_TYPE'].value_counts()
```

```
Out[22]: Business Entity Type 3      42113
          Self-employed            25892
          Other                    9194
          Medicine                 8548
          Government               6958
          Business Entity Type 2    6893
          School                   6316
          Trade: type 7             5608
          Kindergarten             5327
          Construction             4675
          Business Entity Type 1    3830
          Transport: type 4         3540
          Trade: type 3            2584
          Industry: type 9          2403
          Security                 2397
          Industry: type 3          2325
          Housing                  2061
          Bank                     1867
          Police                   1824
          Radiator: type 11         1772
```

## Encode the columns that is required for the model

```
In [23]: # Now We have some categorical column that we will encode foe the model building.
          from sklearn.preprocessing import LabelEncoder
```

```
In [24]: le = LabelEncoder()
```

```
In [25]: loan_df['NAME_CONTRACT_TYPE'] = le.fit_transform(loan_df['NAME_CONTRACT_TYPE'])
          loan_df['CODE_GENDER'] = le.fit_transform(loan_df['CODE_GENDER'])
          loan_df['FLAG_OWN_CAR'] = le.fit_transform(loan_df['FLAG_OWN_CAR'])
          loan_df['FLAG_OWN_REALTY'] = le.fit_transform(loan_df['FLAG_OWN_REALTY'])
          loan_df['NAME_TYPE_SUITE'] = le.fit_transform(loan_df['NAME_TYPE_SUITE'])
          loan_df['NAME_INCOME_TYPE'] = le.fit_transform(loan_df['NAME_INCOME_TYPE'])
          loan_df['NAME_EDUCATION_TYPE'] = le.fit_transform(loan_df['NAME_EDUCATION_TYPE'])
          loan_df['NAME_FAMILY_STATUS'] = le.fit_transform(loan_df['NAME_FAMILY_STATUS'])
          loan_df['NAME_HOUSING_TYPE'] = le.fit_transform(loan_df['NAME_HOUSING_TYPE'])
          loan_df['OCCUPATION_TYPE'] = le.fit_transform(loan_df['OCCUPATION_TYPE'])
          loan_df['WEEKDAY_APPR_PROCESS_START'] = le.fit_transform(loan_df['WEEKDAY_APPR_PROCESS_START'])
          loan_df['ORGANIZATION_TYPE'] = le.fit_transform(loan_df['ORGANIZATION_TYPE'])
```

```
In [26]: loan_df['OCCUPATION_TYPE'].value_counts()
```

```
Out[26]: 8      43437
        14      24254
        3      22669
        10     17787
        4      14546
        6       9264
        0       8135
        11      7223
        16      5218
        2       4633
        1       3637
        12      1975
        9       1443
        15      1047
        17      1010
        13        569
        5        460
        7         425
        Name: OCCUPATION_TYPE, dtype: int64
```

### Print percentage of default to payer of the dataset for the TARGET column

```
In [27]: loan_df.TARGET.value_counts()
```

```
Out[27]: 0      153525
        1      14207
        Name: TARGET, dtype: int64
```

```
In [28]: ((loan_df['TARGET']==1).sum() / loan_df.shape[0])*100
```

```
Out[28]: 8.470059380440226
```

From above code its clear that dataset is highly imbalanced, so we will balanced the dataset.

### Balance the dataset if the data is imbalanced

```
In [29]: # Lets apply the smote oversampling technique
        from imblearn.over_sampling import SMOTE
```

```
In [35]: x = loan_df.drop(['TARGET'], axis=1)
        y = loan_df[['TARGET']]
```

```
In [36]: smt = SMOTE()
```

```
In [37]: x_sm, y_sm = smt.fit_resample(x, y)
```

```
In [39]: print(y_sm.shape)
        print(y_sm.value_counts())
```

```
(307050, 1)
TARGET
0      153525
1      153525
dtype: int64
```

Now dataset is balanced.

```
In [40]: loan_df = pd.concat([x_sm,y_sm], axis=1)
```

```
In [50]: loan_df.head()
```

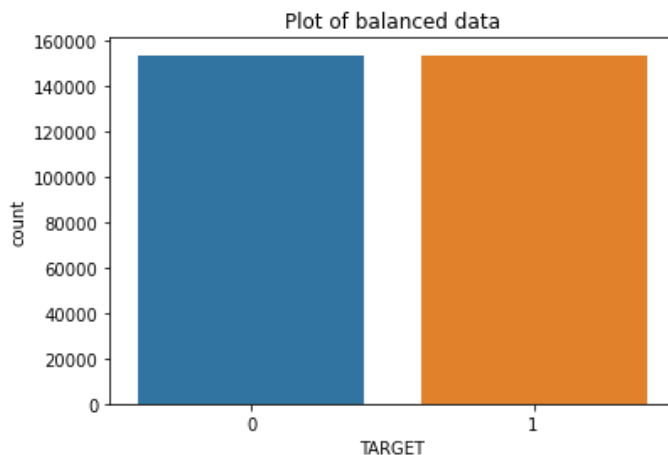
```
Out[50]:
```

	NAME_CONTRACT_TYPE	CODE_GENDER	FLAG_OWN_CAR	FLAG_OWN_REALTY	CNT_CHILDREN	AMT_INCOME_
0	0	1	0	1	0	20
1	1	1	1	1	0	6
2	0	1	0	1	0	9
3	0	0	1	1	1	17
4	0	1	1	1	0	36

5 rows × 72 columns

## Plot the balanced data or imbalanced data

```
In [47]: sns.countplot(loan_df['TARGET'])
plt.title('Plot of balanced data')
plt.show()
```



```
In [51]: x = loan_df.drop(["TARGET"], axis = 1)
y = loan_df[["TARGET"]]
```

```
In [49]: # Now Lets split the data set
from sklearn.model_selection import train_test_split
```

```
In [52]: x_train, x_test, y_train, y_test = train_test_split(x,y, test_size=.25, random_state=10)
```

```
In [53]: from sklearn.preprocessing import StandardScaler
```

```
In [54]: sc = StandardScaler()
```

```
In [56]: x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
```

```
In [58]: x_train.shape
```

```
Out[58]: (230287, 71)
```

```
In [48]: # Now Lets create the ANN model for this problem
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.layers import LeakyReLU, PReLU, ELU, ReLU
from tensorflow.keras.layers import Dropout
```

```
In [57]: classifier = Sequential()
```

```
In [59]: # Lets First introducing the input Layer
classifier.add(Dense(units=71, activation='relu'))
```

```
In [60]: # now Lets introduce the first hidden Layer
classifier.add(Dense(units=50, activation='relu'))
```

```
In [61]: # now Lets introduce the second hidden Layer
classifier.add(Dense(units=35, activation='relu'))
```

```
In [62]: # Adding the output Layer
classifier.add(Dense(1, activation='sigmoid'))
```

```
In [63]: import tensorflow
opt = tensorflow.keras.optimizers.Adam(learning_rate = 0.01)
```

```
In [64]: classifier.compile(optimizer=opt, loss='binary_crossentropy', metrics=['accuracy'])
```

```
In [65]: # early stoping
import tensorflow as tf
early_stopping = tf.keras.callbacks.EarlyStopping(
    monitor = 'val_loss',
    min_delta=0.0001,
    patience=20,
    verbose=1,
    mode='auto',
    baseline=None,
    restore_best_weights=False)
```


```
In [66]: model = classifier.fit(x_train, y_train, validation_split=0.33, batch_size=10, epochs=1000, call

Epoch 1/1000
15430/15430 [=====] - 144s 9ms/step - loss: 0.3602 - accuracy: 0.8473
- val_loss: 0.3223 - val_accuracy: 0.8640
Epoch 2/1000
15430/15430 [=====] - 181s 12ms/step - loss: 0.3084 - accuracy: 0.872
8 - val_loss: 0.2987 - val_accuracy: 0.8775
Epoch 3/1000
15430/15430 [=====] - 235s 15ms/step - loss: 0.2902 - accuracy: 0.883
1 - val_loss: 0.2872 - val_accuracy: 0.8831
Epoch 4/1000
15430/15430 [=====] - 167s 11ms/step - loss: 0.2766 - accuracy: 0.889
5 - val_loss: 0.2775 - val_accuracy: 0.8881
Epoch 5/1000
15430/15430 [=====] - 120s 8ms/step - loss: 0.2761 - accuracy: 0.8933
- val_loss: 0.2748 - val_accuracy: 0.8896
Epoch 6/1000
15430/15430 [=====] - 133s 9ms/step - loss: 0.2656 - accuracy: 0.8954
- val_loss: 0.2722 - val_accuracy: 0.8917
Epoch 7/1000
15430/15430 [=====] - 134s 9ms/step - loss: 0.2648 - accuracy: 0.8967
- val_loss: 0.2808 - val_accuracy: 0.8897
Epoch 8/1000
15430/15430 [=====] - 119s 8ms/step - loss: 0.2610 - accuracy: 0.8986
- val_loss: 0.2760 - val_accuracy: 0.8938
Epoch 9/1000
15430/15430 [=====] - 116s 8ms/step - loss: 0.2608 - accuracy: 0.8992
- val_loss: 0.2873 - val_accuracy: 0.8848
Epoch 10/1000
15430/15430 [=====] - 124s 8ms/step - loss: 0.2605 - accuracy: 0.8997
- val_loss: 0.2631 - val_accuracy: 0.8976
Epoch 11/1000
15430/15430 [=====] - 147s 10ms/step - loss: 0.2618 - accuracy: 0.900
0 - val_loss: 0.2704 - val_accuracy: 0.8960
Epoch 12/1000
15430/15430 [=====] - 146s 9ms/step - loss: 0.2621 - accuracy: 0.9003
- val_loss: 0.2954 - val_accuracy: 0.8952
Epoch 13/1000
15430/15430 [=====] - 147s 10ms/step - loss: 0.2581 - accuracy: 0.901
5 - val_loss: 0.2639 - val_accuracy: 0.8968
Epoch 14/1000
15430/15430 [=====] - 151s 10ms/step - loss: 0.2592 - accuracy: 0.902
4 - val_loss: 0.2715 - val_accuracy: 0.8936
Epoch 15/1000
15430/15430 [=====] - 167s 11ms/step - loss: 0.2570 - accuracy: 0.903
3 - val_loss: 0.2760 - val_accuracy: 0.8925
Epoch 16/1000
15430/15430 [=====] - 152s 10ms/step - loss: 0.2551 - accuracy: 0.903
2 - val_loss: 0.2673 - val_accuracy: 0.8964
Epoch 17/1000
15430/15430 [=====] - 172s 11ms/step - loss: 0.2559 - accuracy: 0.903
8 - val_loss: 0.2720 - val_accuracy: 0.8979
Epoch 18/1000
15430/15430 [=====] - 153s 10ms/step - loss: 0.2592 - accuracy: 0.903
0 - val_loss: 0.3138 - val_accuracy: 0.8935
Epoch 19/1000
15430/15430 [=====] - 153s 10ms/step - loss: 0.2573 - accuracy: 0.903
2 - val_loss: 0.2728 - val_accuracy: 0.8968
Epoch 20/1000
15430/15430 [=====] - 191s 12ms/step - loss: 0.2567 - accuracy: 0.903
6 - val_loss: 0.2710 - val_accuracy: 0.8964
Epoch 21/1000
15430/15430 [=====] - 156s 10ms/step - loss: 0.2614 - accuracy: 0.902
7 - val_loss: 0.2786 - val_accuracy: 0.8967
Epoch 22/1000
15430/15430 [=====] - 169s 11ms/step - loss: 0.2563 - accuracy: 0.904
8 - val_loss: 0.2691 - val_accuracy: 0.8962
Epoch 23/1000
15430/15430 [=====] - 166s 11ms/step - loss: 0.2570 - accuracy: 0.904
5 - val_loss: 0.2726 - val_accuracy: 0.8940
Epoch 24/1000
```



```
15430/15430 [=====] - 153s 10ms/step - loss: 0.2585 - accuracy: 0.904
1 - val_loss: 0.2848 - val_accuracy: 0.8906
Epoch 25/1000
15430/15430 [=====] - 167s 11ms/step - loss: 0.2538 - accuracy: 0.906
0 - val_loss: 0.2698 - val_accuracy: 0.8997
Epoch 26/1000
15430/15430 [=====] - 153s 10ms/step - loss: 0.2541 - accuracy: 0.905
9 - val_loss: 0.2792 - val_accuracy: 0.8971
Epoch 27/1000
15430/15430 [=====] - 148s 10ms/step - loss: 0.2540 - accuracy: 0.905
8 - val_loss: 0.2617 - val_accuracy: 0.9019
Epoch 28/1000
15430/15430 [=====] - 149s 10ms/step - loss: 0.2538 - accuracy: 0.905
7 - val_loss: 0.2621 - val_accuracy: 0.8995
Epoch 29/1000
15430/15430 [=====] - 131s 8ms/step - loss: 0.2566 - accuracy: 0.9047
- val_loss: 0.2733 - val_accuracy: 0.8953
Epoch 30/1000
15430/15430 [=====] - 141s 9ms/step - loss: 0.2581 - accuracy: 0.9048
- val_loss: 0.2685 - val_accuracy: 0.9002
Epoch 31/1000
15430/15430 [=====] - 140s 9ms/step - loss: 0.2561 - accuracy: 0.9063
- val_loss: 0.2729 - val_accuracy: 0.8966
Epoch 32/1000
15430/15430 [=====] - 131s 8ms/step - loss: 0.2538 - accuracy: 0.9067
- val_loss: 0.2778 - val_accuracy: 0.8981
Epoch 33/1000
15430/15430 [=====] - 133s 9ms/step - loss: 0.2556 - accuracy: 0.9060
- val_loss: 0.2680 - val_accuracy: 0.8994
Epoch 34/1000
15430/15430 [=====] - 145s 9ms/step - loss: 0.2561 - accuracy: 0.9058
- val_loss: 0.2655 - val_accuracy: 0.9002
Epoch 35/1000
15430/15430 [=====] - 154s 10ms/step - loss: 0.2542 - accuracy: 0.906
0 - val_loss: 0.2722 - val_accuracy: 0.8973
Epoch 36/1000
15430/15430 [=====] - 146s 9ms/step - loss: 0.2552 - accuracy: 0.9062
- val_loss: 0.2675 - val_accuracy: 0.8985
Epoch 37/1000
15430/15430 [=====] - 141s 9ms/step - loss: 0.2520 - accuracy: 0.9077
- val_loss: 0.2634 - val_accuracy: 0.9010
Epoch 38/1000
15430/15430 [=====] - 143s 9ms/step - loss: 0.2567 - accuracy: 0.9063
- val_loss: 0.2817 - val_accuracy: 0.8927
Epoch 39/1000
15430/15430 [=====] - 144s 9ms/step - loss: 0.2585 - accuracy: 0.9055
- val_loss: 0.2759 - val_accuracy: 0.8945
Epoch 40/1000
15430/15430 [=====] - 147s 10ms/step - loss: 0.2537 - accuracy: 0.907
6 - val_loss: 0.2709 - val_accuracy: 0.9000
Epoch 41/1000
15430/15430 [=====] - 148s 10ms/step - loss: 0.2508 - accuracy: 0.908
7 - val_loss: 0.2691 - val_accuracy: 0.9010
Epoch 42/1000
15430/15430 [=====] - 135s 9ms/step - loss: 0.2516 - accuracy: 0.9085
- val_loss: 0.2692 - val_accuracy: 0.9007
Epoch 43/1000
15430/15430 [=====] - 103s 7ms/step - loss: 0.2502 - accuracy: 0.9092
- val_loss: 0.2662 - val_accuracy: 0.9008
Epoch 44/1000
15430/15430 [=====] - 103s 7ms/step - loss: 0.2519 - accuracy: 0.9089
- val_loss: 0.2788 - val_accuracy: 0.8976
Epoch 45/1000
15430/15430 [=====] - 104s 7ms/step - loss: 0.2546 - accuracy: 0.9078
- val_loss: 0.2718 - val_accuracy: 0.8973
Epoch 46/1000
15430/15430 [=====] - 103s 7ms/step - loss: 0.2536 - accuracy: 0.9075
- val_loss: 0.2758 - val_accuracy: 0.8983
Epoch 47/1000
15430/15430 [=====] - 103s 7ms/step - loss: 0.2547 - accuracy: 0.9078
- val_loss: 0.2637 - val_accuracy: 0.8997
Epoch 47: early stopping
```

```
In [67]: # Summarize the history for accuracy
plt.plot(model.history['accuracy'])
plt.plot(model.history['val_accuracy'])
plt.title('model_accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
plt.show()
```



```
In [80]: # Making the prediction and evaluating the model.
y_pred = classifier.predict(x_test)
y_preds = (y_pred>=0.5)

2399/2399 [=====] - 13s 5ms/step
```

```
In [81]: y_preds
```

```
Out[81]: array([[ True],
 [ True],
 [False],
 ...,
 [False],
 [ True],
 [ True]])
```

```
In [82]: from sklearn.metrics import classification_report, confusion_matrix, auc, roc_curve, roc_auc_score
```

```
In [83]: print(classification_report(y_test, y_preds))
```

	precision	recall	f1-score	support
0	0.87	0.94	0.90	38389
1	0.94	0.86	0.90	38374
accuracy			0.90	76763
macro avg	0.90	0.90	0.90	76763
weighted avg	0.90	0.90	0.90	76763

### Calculate Sensitivity as a metric

```
In [84]: confusion = confusion_matrix(y_test, y_preds)
```

```
In [85]: confusion
```

```
Out[85]: array([[36222, 2167],
 [ 5484, 32890]], dtype=int64)
```

```
In [86]: # True Positive
TP = confusion[1,1]
# True Negatives
TN = confusion[0,0]
# False Positives
FP = confusion[0,1]
# False Negatives
FN = confusion[1,0]
```


```
In [87]: sensitivity = TP/(TP+FN)
print("Sensitivity: %.3f" %sensitivity)
```

Sensitivity: 0.857

## Calculate area under receiver operating characteristics curve

```
In [90]: # Let's compute the AUC(ROC) Receiver Operating Characteristics curve and also visualize by th
probs = classifier.predict(x_test)
# Area under the curve.
auc_cl = roc_auc_score(y_test, probs)
print("AUC: %.3f" %auc_cl)
# Calculating the roc curve.
fpr, tpr, thresholds = roc_curve(y_test, probs)
#plotting the auc(roc) curve.
plt.plot([0,1], [0,1], linestyle='--')
plt.plot(fpr,tpr, marker='.')
plt.title("AUC(ROC) Curve")
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.show()
```

2399/2399 [=====] - 13s 5ms/step  
AUC: 0.939



In [ ]:

In [ ]:

In [ ]: